



What You Need to Know About the Serverless Edge

Who am I?

- Developer Experience Engineer at [LaunchDarkly](#)
- Co-author of [The Jamstack Book from Manning](#)
- Editor of the [Jamstacked newsletter from Cooper Press](#)
- Run the [CFE.dev community](#) and Orlando Devs

LaunchDarkly ➔

What is "the edge"?



The Edge's real name is David Howell Evans. He was born in August 1961 and is the lead guitarist, keyboardist, and backing vocalist of U2.

Questions?

- Mastodon: [@remotesynth@mastodon.xyz](https://mastodon.xyz/@remotesynth)
- Email: brinaldi@launchdarkly.com

Ok. I know this ran a little short,
but any questions?

A close-up photograph of a man with dark hair and a mustache, wearing a camouflage jacket. He has a weary or concerned expression, with his eyes half-closed and a slight frown. The background is dark and out of focus.

When enough people make false promises

gifs.com

In all seriousness though, before we discuss what the edge is, I want to talk about broken promises. As developers, I think we've all gotten used to being sold a new technology only to find out later that it doesn't quite live up to the hype.

FEATURE

Edge Computing: What It is and How It's a Game-Changer

5 MINUTE READ

ta / 19 Oct 2022 / AR/VR / Cloud / ReadWrite

edge functions have been such a game changer, like the ability to run arbitrary code over an http request w like basically no overhead, possibilities r w/o end

 **Larry Hudson** @larryhudsondev · May 28

I recently got an Apple Watch, and I've been really impressed with what you can do with Apple's Shortcuts app (across Mac, iOS and Watch).

If you're comfortable with JavaScript and data APIs, a great match for Shortcuts is @Netlify Edge Functions.

[Show this thread](#)

10:58 AM · May 29, 2022 · Twitter for iPhone

Why Serverless will enable the Edge Computing Revolution

Serverless changes how we build applications both technically and conceptually — the same changes will enable Edge Computing adoption.

Cloud at the Edge: A Major Game-Changer Waiting on the Wings

MARKETS

Edge Computing Is Leading The Next Great Tech Revolution



And let's be honest, edge computing and edge serverless specifically have gotten a ton of hype lately. In many cases, the hype is coming from within companies that have a stake in the game.



About ▾ Blog Newsletter Podcasts ▾ Products Resources Sponsorships

Lower My AWS Bill

11.10.2021

The Unfulfilled Promise of Serverless

BY COREY QUINN

I suggest that serverless computing, or “serverless” has hype that at this point has outpaced what the technology / philosophy / religion has been promising. Serverless computing arrived (debatably; please do not...



Prev

[Home](#) ▶ [Blog](#) ▶ The Unfulfilled Promise of Serverless

Next

I suggest that serverless computing, or “serverless” has hype that at this point has outpaced what the technology / philosophy / religion has been promising. Serverless computing arrived (debatably; please do not email me about this whatever you do, fans of Google App Engine / CGI scripts / managed SaaS offerings /

And serverless arguably has a history of broken promises. Corey Quinn, a very well known expert on AWS and serverless, wrote this in November 2021. Now I'm a big believer in serverless in general, but it's hard to disagree that in many cases it hasn't lived up to the hype. <https://www.lastweekinaws.com/blog/the-unfulfilled-promise-of-serverless/>

CHRIS COYIER

Bio Timeline Interviews Speaking
Photos Random RSS

MAY 4, 2022

It doesn't much matter how CDNy your Jamstack site is if everything important happens from a single-origin server. Edge functions are probably part of the solution.

By "everything important", I mean hitting APIs. It's not that I don't think it's a neat idea to pre-build and CDN-serve an HTML shell page and then make

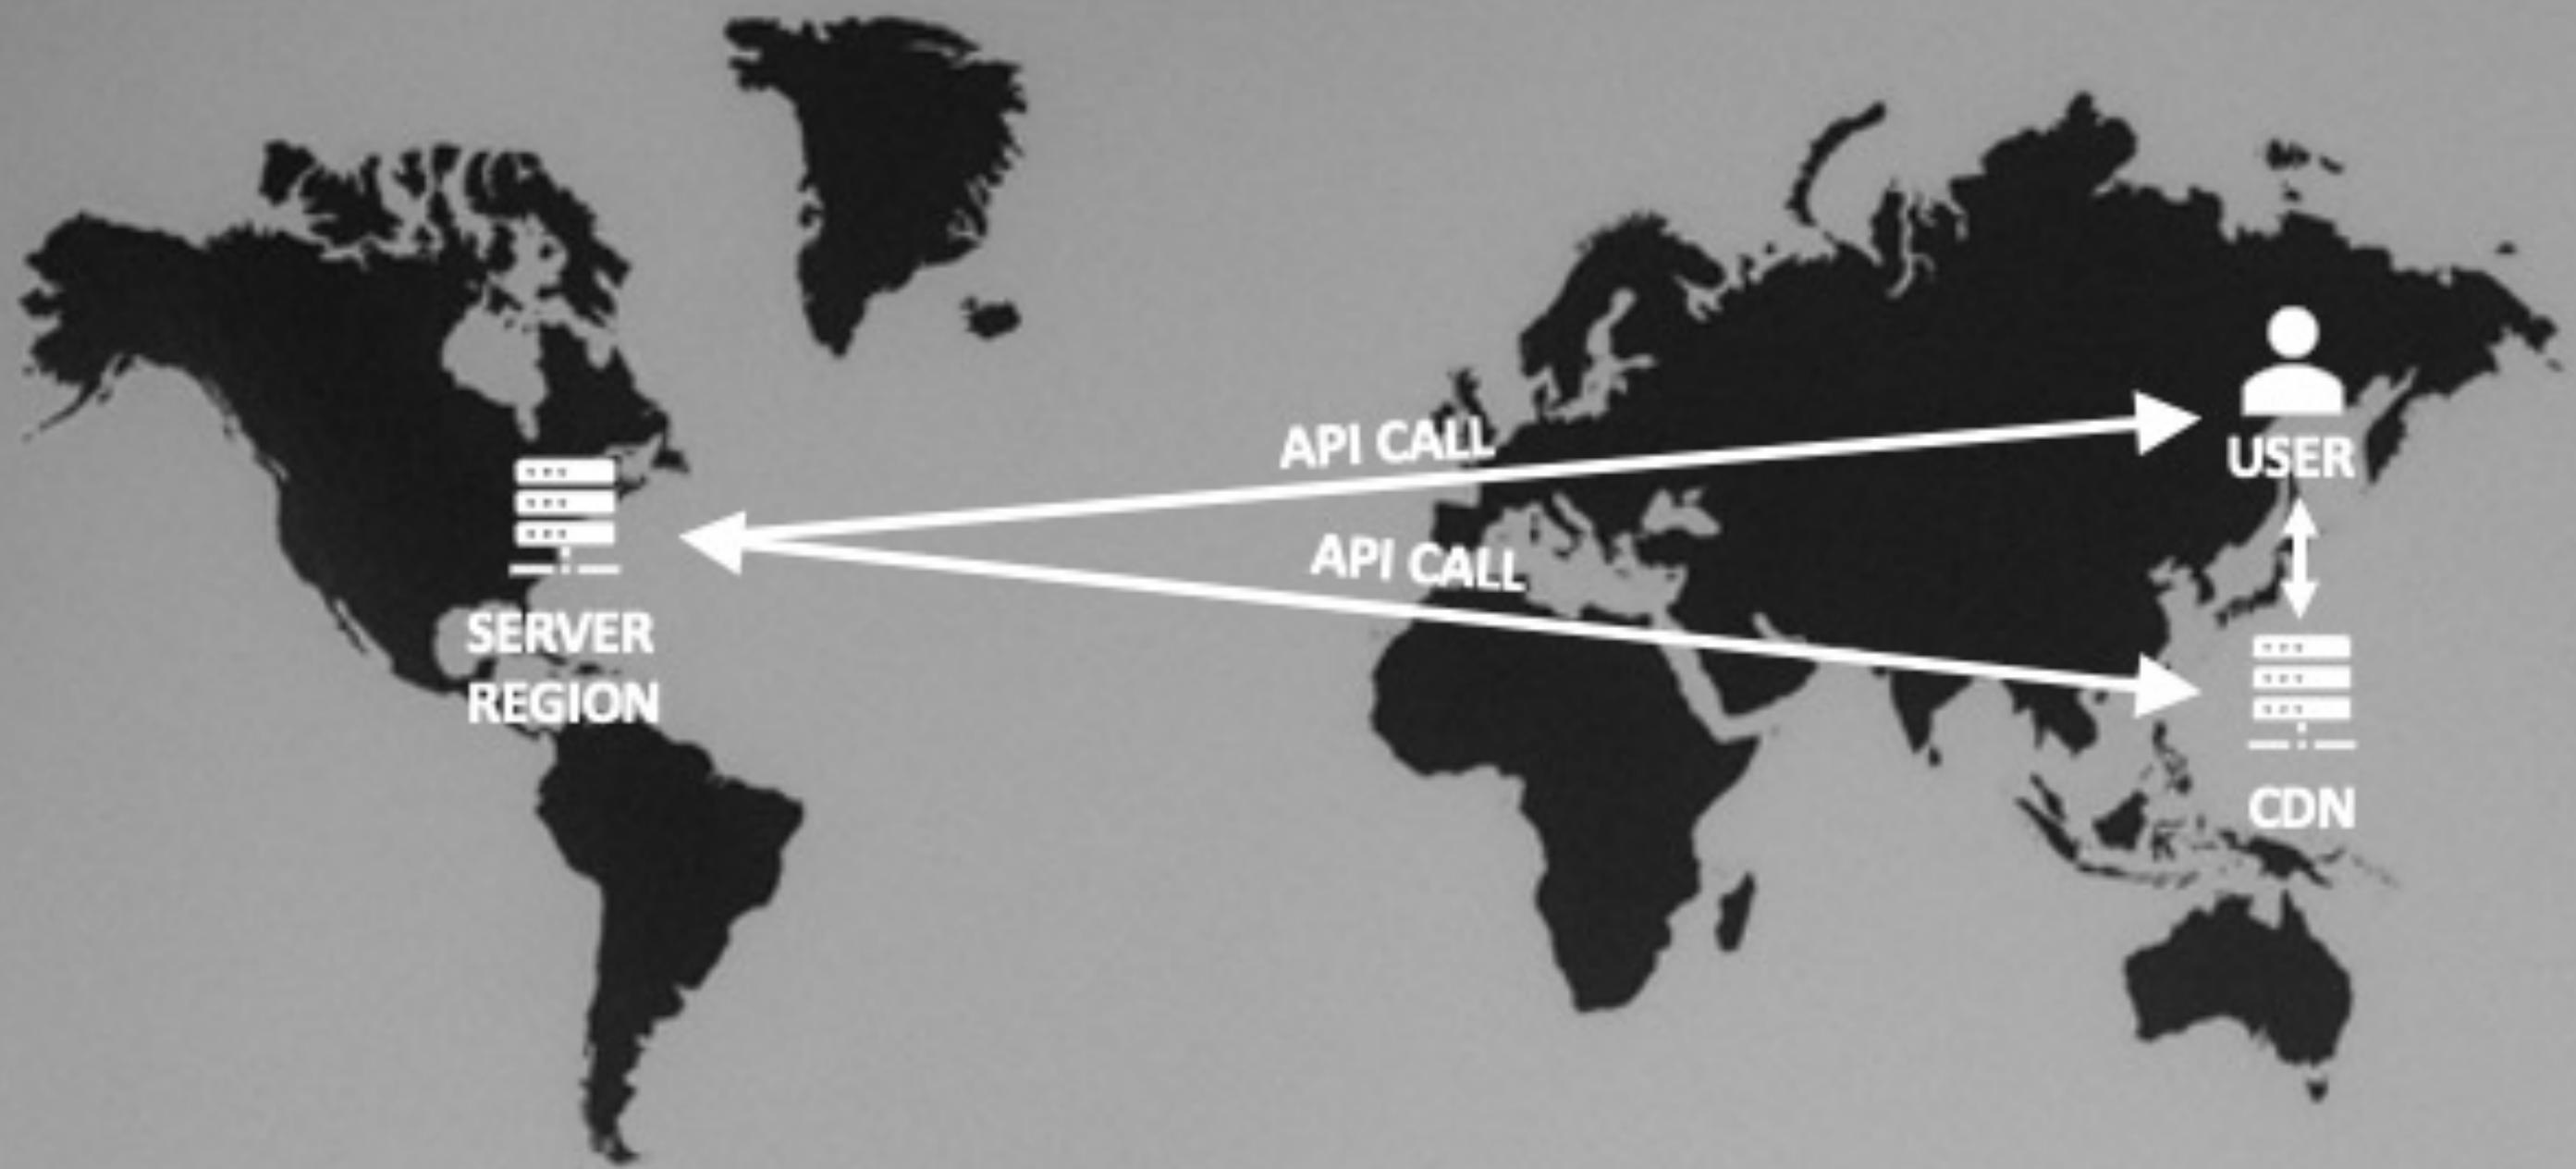
And the edge in particular already has some baggage. As you hopefully could tell from the intro, I am also a firm believer in the Jamstack architecture for building sites, which deployed site assets to the edge. But as Chris Coyier points out in this post from May, you probably still need data and that data is still coming from a specific location or region. <https://chriscoyier.net/2022/05/04/it-doesnt-much-matter-how-cdny-your-jamstack-site-is-if-everything-important-happens-from-a-single-origin-server-edge-functions-are-probably-part-of-the-solution/>

Name	Code
US East (Ohio)	us-east-2
US East (N. Virginia)	us-east-1
US West (N. California)	us-west-1
US West (Oregon)	us-west-2
US GovCloud West	us-gov-west-1
US GovCloud East	us-gov-east-1
Asia Pacific (Hong Kong)	ap-east-1
Asia Pacific (Mumbai)	ap-south-1
Asia Pacific (Seoul)	ap-northeast-2
Asia Pacific (Singapore)	ap-southeast-1
Asia Pacific (Sydney)	ap-southeast-2
Asia Pacific (Tokyo)	ap-northeast-1
Canada (Central)	ca-central-1
Europe (Frankfurt)	eu-central-1
Europe (Ireland)	eu-west-1
Europe (London)	eu-west-2
Europe (Paris)	eu-west-3
Europe (Stockholm)	eu-north-1
Middle East (Bahrain)	me-south-1
South America (São Paulo)	sa-east-1
China (Beijing)	cn-north-1
China (Ningxia)	cn-northwest-1

For example, if you're on AWS you'll deploy your lambda function to one of their 31 different regions across the globe.



In essence we moved from the traditional deployment model that looked like this, where the user, regardless of their location, was served a site from a single location or region.



To a model like this, where the user would be served static assets from a region close to them but then all of the data which came from APIs, data stores or databases was served from a single region. This means that the static assets load fast, but similar latency still exists populating the page data.



Edge functions aim to solve this problem (and not just for Jamstack sites), by bringing some - or all - of the backend processing up to the CDN level in proximity to the user. So it's no longer just static assets on the CDN but actual compute is being performed there as well.




CDN Node



User



Origin Server



In an ideal world, it might look something like this image taken from the Remix blog, where every user request has static assets fulfilled by the CDN and backend assets also fulfilled by the CDN via edge functions.

Promise: Reduced latency

-  Backend calls go to the server geographically closest to the user.
-  Geographic latency may not be the critical issue in your site's performance

So that's the biggest selling point of edge functions: that they will improve the performance of your app by reducing the latency. And they do actually achieve that promise in terms of the actual function call. However, if your edge function needs data, that data may still be region locked. And anyway, unless you are building games or IoT, that latency caused by distance, while not worth ignoring, may not be the critical aspect to improving your site's performance.

Just How Much Does Latency Affect Your Speed?

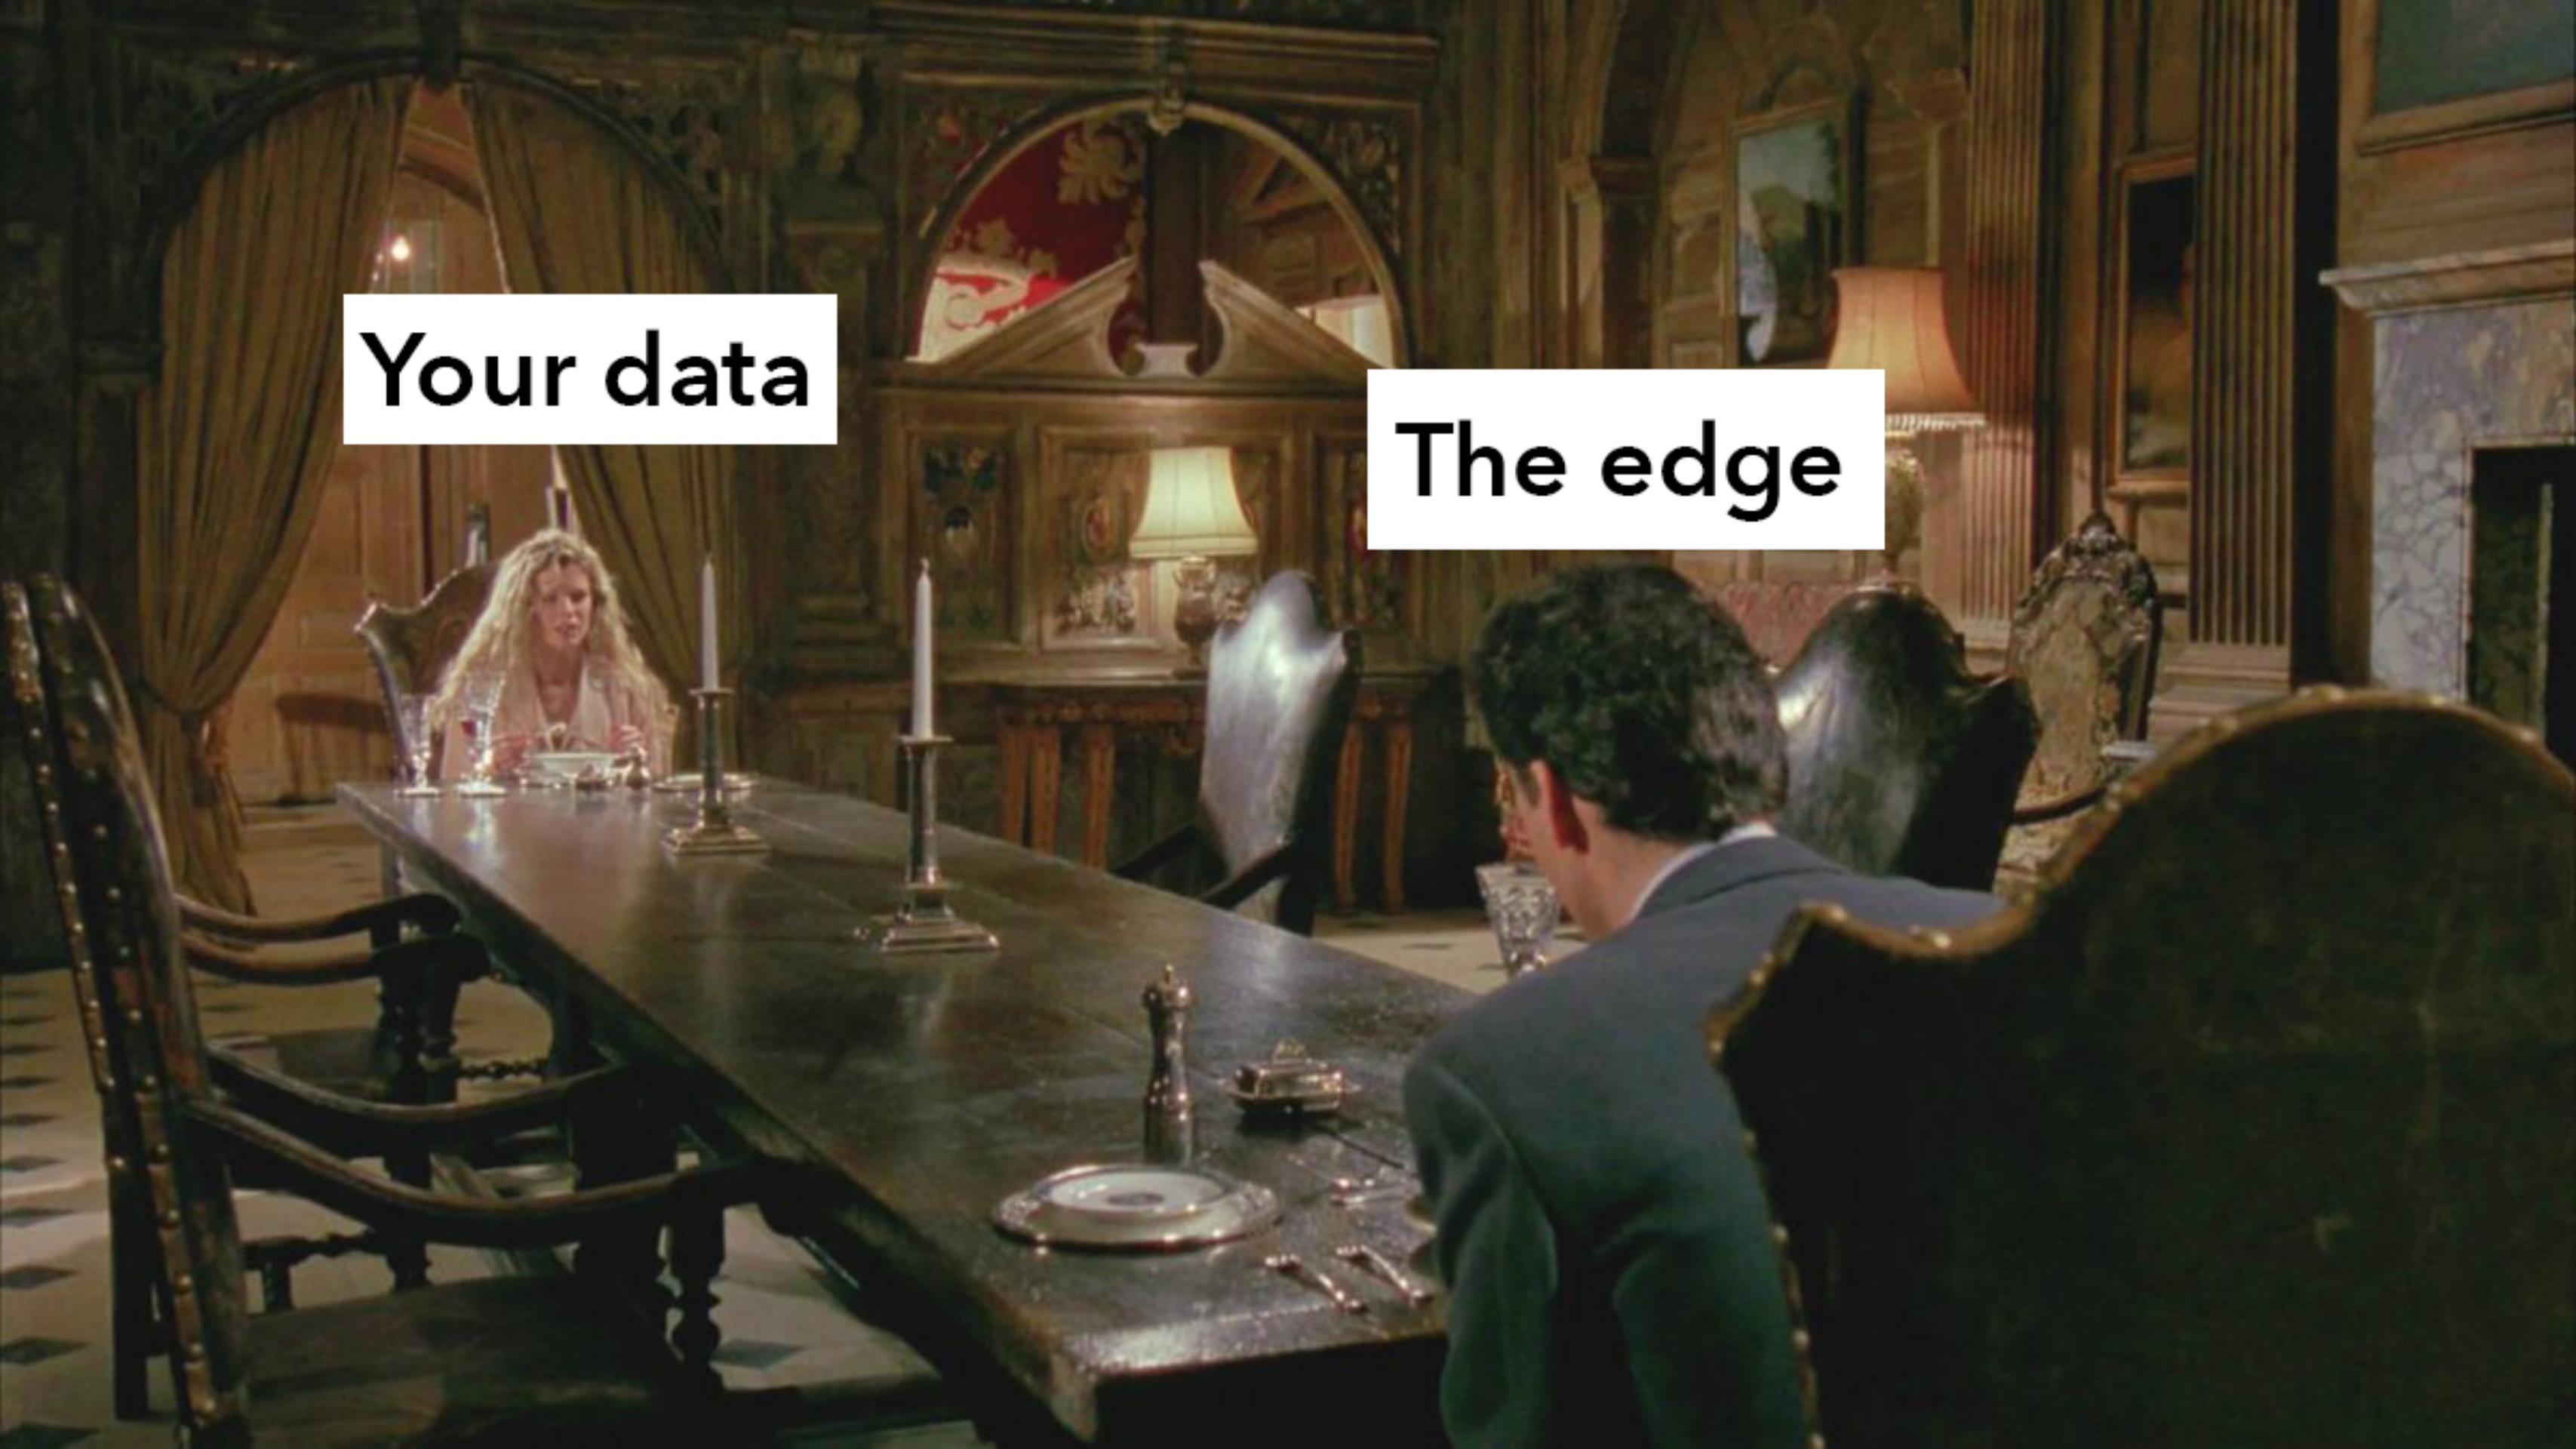
Well, let's consider the perfect example of having a high-speed fiber-optic network to every point in the globe and a speed of light that is approximately 200,000 km/second.

At that insanely fast speed, it will take around **150ms to reach from the US West Coast to central Asia**. This theoretical scenario assumes that every point on the internet is a fiber-optic network, which we know is not the case for every location in the world, and ignores the impact of routing devices and the numerous hops that traffic will need to make along the way. **In reality, such a trip across the globe would take several seconds to complete.**

However, even if we assume the fastest scenario of 150ms to connect to a website across the globe, a user will start to notice the impact. As data needs to move both there and back, it will mean around a 300ms delay in data getting back to the user.

300ms might not sound like a noticeable delay when surfing a website if we think of it as a one-off delay. But we are not sending data only once. The reality is that there is a constant stream of messages going back and forth to provide relevant content and to ensure the reliability and security of the network. That 300ms delay affects every one of those communications.

I'm not trying to dismiss the cost of geographic latency because it definitely can add up. As this article points out, data going from US West to Central Asia, for example, could take up to a few seconds. So this is a legitimate selling point that edge functions can live up to when used properly. <https://www.snapt.net/blog/how-geographic-distance-affects-latency>



Your data

The edge

What do I mean by "used properly"? Well that gets at the biggest potential drawback to edge functions when it comes to reducing latency.

Promise: Reduced latency

-  Backend calls go to the server geographically closest to the user.
-  Geographic latency may not be the critical issue in your site's performance
-  Your database or other backend APIs may still be region locked

And that's that your data in the form of APIs, third-party services or databases may still exist in specific region. There are solutions that can help mitigate this such as edge enabled databases or Cloudflare's new smart placement but the key thing to remember is that, if you're going to access data in your edge function, you need to have a plan or you may not see the performance improvement from reduced latency you'd hoped for.

Promise: It's easy

-  It's just JavaScript, so there's a low barrier to entry and lots of framework support
-  Runtimes are often not Node-based, so you may not have all your favorite tools
-  Every provider is unique and choosing one isn't easy

Another promise you may hear about edge functions is that it's super easy to get started. And there's truth to this. In every case you can just use JavaScript, though some providers do support other languages as well. And lots of frameworks are available to help you build them. But the runtimes are mostly not Node-based, so you can't always bring your npm-based tools with you. And each provider implements them differently so it can be tough to choose and difficult to learn their offering and quirks. And your code is generally not easily portable.

Promise: Additional capabilities

-  Edge functions can intercept the request or response
-  Um...actually, this is just cool
-  I got nuthin

Where I get more excited about edge functions from a developer standpoint is where it can actually allow me to do things that previously were much more difficult. And that is where the ability to intercept the request and response comes in.




Essentially, your edge function - or middleware as some frameworks call it in this scenario - will be called when the request comes in before it hits the origin server. It can modify that request, redirect it - so no need for slower server-side or client-side redirects. You can even handle access and authentication here.



Or it can even fulfill the response itself. For example, if you have static assets like HTML, CSS, images and more cached on the edge, you can simply respond without hitting the origin server. This is actually how Cloudflare Pages, which we'll talk about, works.



If you haven't redirected or responded directly, the request continues to the origin server to be fulfilled.




When that response comes back, your edge function can intercept it again. In this case, you can do things like modify the response. For example, you might add headers that have user specific information or you may even modify the HTML of the response.

Add personalization to static HTML with Netlify Edge Functions — no browser JavaScript required

Personalization — the process of creating customized experiences for visitors to a website — is a hot topic on the web in 2022. In a crowded world of digital experiences clamoring for our attention, businesses are turning to personalization to connect more deeply with their customers, and ultimately, generate more sales and growth.

And shipping less client-side JavaScript is an even hotter topic. New and emerging front-end frameworks such as [Astro](#) and [Eleventy](#) are pushing the boundaries of the JavaScript framework era of the late 2010s, making it their core mission to ship less and less JavaScript *by default*, and bring developers' focus back to a lean, native and performant web.

The great news is that with [Netlify Edge Functions](#) you can achieve dynamic personalization with *no* client-side JavaScript — resulting in a great developer



by [Salma Alam-Naylor](#)

💡 I write code for your entertainment

 [Watch me live](#)

📅 11 Aug 2022

⌚ 8 min read

▼ Table of contents

[Deploy the tutorial code to Netlify](#)

This enables you to do some really cool things such as make a dynamic, user-specific response based upon a completely static HTML asset. Salma Alam-Naylor has a number of examples doing just that using Netlify's edge functions. <https://whitep4nth3r.com/blog/add-personalization-to-static-html-with-edge-functions-no-browser-javascript/>

Use Cases

- A/B Testing
- Updating/Modifying/Custom HTTP headers
- Conditional routing
- SEO
- User authentication/authorization
- IoT/Gaming
- Personalization
- Compliance


With all that in mind, here are some of the things you might consider doing using edge functions. You can replace out content for A/B tests without any flash of rendering. You can add custom http headers. You can route user's to a new location. You can serve a different version of your site for search crawlers. You can check for authentication and redirect otherwise. Remove critical latency on games or IoT apps. Personalize the output, for example show login state without flash of rendering. Serve up different versions depending on the user location to meet location specific compliance issues.

Side note: Limitations

Most services have some degree of limitations for edge functions including limits on:

- Code size
- CPU execution time
- Memory usage
- Response time

One quick side note is that almost every service has much stricter limitations on edge functions than a standard serverless functions. These tend to differ between providers, so it's worth researching before you begin a project that might hit them, but also means not everything can move to the edge.



Hopefully that gave you the gist of what edge functions are and where they can be especially useful. Let's move on to the various options for implementing them.

Core Edge Serverless

- Cloudflare Workers
- Deno Deploy
- AWS (Lambda@Edge, Cloudfront Functions)
- Fastly Compute@Edge
- Akamai EdgeWorkers
- Stackpath
- Azion

When I think about edge functions there is basically two different types of offerings. The first I'm calling core edge serverless. These are largely CDN companies who have added compute offerings to their products. This isn't a complete list but some of the more well known players.

Edge Platforms

- [Netlify Edge Functions \(Deno\)](#)
- [Vercel Edge Functions \(Cloudflare\)](#)
- [Supabase Edge Functions \(Deno\)](#)
- [Cloudflare Pages Functions \(Workers\)](#)

Then there are what I am calling edge platforms. These are edge function offerings that are built on top of those core edge serverless offerings but also provide improved developer experience and integrated continuous deployment.

Framework Support

- [Next.js](#)
- [Astro \(Example the Netlify SSR Integration\)](#)
- [SvelteKit \(See examples from Vercel and Netlify\)](#)
- [Eleventy \(Using the Eleventy Edge plugin that is Netlify only\)](#)
- [Remix \(Using adapters\)](#)
- [Fresh \(Built by Deno\)](#)

There are also the front-end frameworks that help you build applications using edge functions usually by adding conventions that hide the underlying implementation and complexity from you.

Edge Enabled Databases

- Fauna – A serverless database with built in global replication and strong consistency.
- Planetscale Portals – A serverless MySQL database with support for replicating to read-only regions.
- Cloudflare DI – A SQLite based relational database that has globally replicated reads.
- DGraph – A distributed GraphQL database with a graph backend.
- Upstash – A serverless, replicated Redis data store (runs Vercel's KV offering)

When it comes to moving your data to the edge, there is a fast growing list of offerings that are full databases – rather than just cache or KV stores – that are edge enabled. In most cases, the reads are replicated but the writes are region locked, which is likely sufficient for your needs. It's worth keeping in mind that most of these are eventually consistent, which means there may be a lag between the data being updated and replicated to a given region.

2023 State of Databases for Serverless & Edge

An overview of 30+ solutions and the emerging trends.



LEE ROBINSON
JAN 28, 2023

41 5

Share ...

There's been massive innovation in the database and backend space for developers building applications with serverless and edge compute. There are new tools, companies, and even programming models that simplify how developers store data.

This post will be an overview of databases that pair well with modern application and compute providers.

Criteria

I'll focus on **transactional**¹ workloads instead of analytical² workloads.

The "backend" space is vast: search, analytics, data science, and more – so I'll niche down here. The primary criteria of this overview is:

1. Services that work exceptionally well when paired with serverless and edge compute
2. Services that work with JavaScript and TypeScript codebases

For a good overview of database solutions for the edge, check out this post from Lee Robinson of Vercel. It covers a lot of options, including a ton I haven't discussed here.

Cloudflare Workers



Let's start our more detailed look with Cloudflare Workers. In this area, Cloudflare has arguably one of the most mature products. They also have an extensive network of 275 CDN locations where your function gets replicated to.

Cloudflare Workers

Workers can be deployed via the CLI tool called Wrangler. They run in the Cloudflare Worker Runtime based on V8 that supports JavaScript/Wasm.¹

¹The core Workers runtime was recently open-sourced under the name [Workerd](#) (meaning you can self-host Workers).

Workers use non-Node runtime based upon V8 that allows you to build functions using JavaScript. They also recently open sourced this runtime under the unusual name of Workerd, which allows you to self-host Workers if you choose to.

Cloudflare Workers

A default worker::

```
addEventListener('fetch', event => {
  event.respondWith(handleRequest(event.request));
});

async function handleRequest(request) {
  return new Response('Hello worker!', { status: 200 });
}
```


The fetch event is any incoming HTTP request (ex. *.workers.dev or a custom origin). In most cases you will send a response, and if you are using a workers.dev it's required. This can be a custom response like the one above or the result of a fetch to return resources.

Cloudflare Pages Functions

- Automatically deploy Workers from a Cloudflare Pages site by placing Worker files in the functions directory ([source](#)).
- Billed the same as Workers.

You can also easily deploy workers as part of a project hosted on Cloudflare Pages by simply placing your Workers files in a functions directory.

Cloudflare Pages Functions



This allows for things like dynamic routes and catchall routes for your Worker-based application API. You can also add middleware in any or all of these directories which will run on every call.

Data at the Edge with Cloudflare

1. **KV** - A simple key value store that replicates to the edge. Eventually consistent.
2. **Durable objects** - Object storage (the stored objects are an instance of a class in JavaScript). Objects can have storage attached and exist in one location at any one time, giving them strong transactional consistency. This is available only on paid accounts.
3. **R2** - This is S3 compatible object storage. Available free with a pay-as-you-go option if you exceed the free limits.
4. **DI** - A SQLite-based SQL database with read-only replication to the edge. Currently available in an open alpha.

Cloudflare has a ton of edge data options depending on your needs. They offer everything from a key value store to various degrees of caching and even a edge enabled SQL database that is currently in open alpha.

Cloudflare + LaunchDarkly

Gist: <https://gist.github.com/remotesynth/c308dff55252bbc4e98500eff9d3433a>

Before we move on, I wanted to show one more quick example because it illustrates two things, Cloudflare's HTMLRewriter that makes it easy to modify the response HTML and LaunchDarkly's edge integration for Cloudflare. This example uses the HTMLRewriter combined with LaunchDarkly's edge SDK that works with Cloudflare to modify headers and run A/B tests using a Worker.

Deno Deploy

- Supports 35 Regions
- Added support for a KV using SQLite that is replicated across at least 6 data centers, spanning 3 regions (US, Europe, and Asia). Support for the cache API is on the roadmap
- Support JSX (and TSX) out of the box which can be useful for SSR using the `renderSSR` function (source)
- Also not Node-based but claim support for over a million npm modules.

Deno has 34 regions currently on their edge. One drawback is the lack of support for a cache API on the edge but they recently added data support on the edge and Deno does support outbound TCP and TLS connections to connect to databases like PostgreSQL, SQLite or MongoDB. There's good examples of connecting Deno to edge data services this way. They support JSX out of the box, which can be useful for SSR React apps. Their GitHub integration makes it easy to deploy a Deno app. While Deno is not Node-based, there are a number ported and you can use npm modules that aren't ported provided they use ES Modules format, though this means a lot of existing modules may not work.

Deno Deploy

Gist: <https://gist.github.com/remotesynth/97c69486ab8defd0baa17c238c99cae2>

This is a simple example that shows me calling the LaunchDarkly API to make changes to a simple HTML response. Note that I can't use the existing Node SDK because it won't work in Deno currently. (In the Deno dash <https://dash.deno.com/playground/salty-donkey-32>)

AWS Lambda@Edge and CloudFront Functions

- Lambda@Edge works just like Lambda but is replicated across 13 regional edge caches.
- CloudFront Functions are replicated across 218 edge locations but are currently limited to view triggers and have limited capabilities.
- No edge databases

I want to cover AWS quickly even though it's still the giant when it comes to serverless, but its edge offerings are pretty standard and, to be honest, are a little confusing. Lambda@Edge lets you deploy a regular Lambda to one of its 13 regional edge caches. If you need more edge nodes, CloudFront Functions offer that but you are extremely limited in what you can do. They don't have any specific edge data offerings that I am aware of, so your data is still probably region locked - though they seem to release a new service daily.



An Introduction To AWS CloudFront Functions

CloudFront Functions allow you to deploy serverless JavaScript functions to AWS' network of edge locations to be executed as close as possible to end-users. This article will get you started.

By Ayoluwa Isaiah 

#javascript

Mar 7, 2022

Amazon CloudFront is a service that speeds up the distribution and delivery of static and dynamic web content through its global network of machines spread across hundreds of locations, also known as edge locations. CloudFront Functions are an incarnation of FaaS (Function as a Service) and allow you to deploy JavaScript functions to AWS' network of edge locations, to be executed as close as possible to end users.

This new feature allows you to customize or personalize content for your application users closer to where they're located, thus minimizing network latency. For example, you can transform HTTP headers or API responses to personalize your application for each visitor, implement authentication or encryption logic (such as JWT authentication) to allow or deny requests, or set up URL rewrites and redirects right on the edge.

Honestly, I find the differences between Lambda@Edge and Cloudfront functions difficult to understand but this article was the best explanation I've seen. Basically Cloudfront functions are great for things like programmatic redirects or auth. Most everything else you'll use Lambda@Edge.
<https://www.honeybadger.io/blog/aws-cloudfront-functions/>

Netlify Edge Functions

- Add a `netlify/edge-functions` directory at the root of your project.
- Edge functions provide geolocation information in the `context.geo` variable.
- No support (yet) for the Deno KV
- Netlify also has a Next.js Advanced Middleware package that includes things like modifying headers and content.
[source](#).

Netlify makes deploying edge functions to Deno for your project pretty easy and integrated with your project code. The easiest way is to just add a `netlify/edge-functions` directory to your project, though this is configurable. As with almost all edge offerings, edge functions automatically provide geolocation data. They also have an advanced middleware offering that does some cool stuff for Next.js to make it easier to modify the source. Otherwise, it's the same as Deno in that you may run into limitations with some npm modules but note that there isn't support for the Deno KV.

Netlify Edge Functions

```
import { Context } from "https://edge.netlify.com";
import { DOMParser } from "https://esm.sh/linkedom";

export default async (request: Request, context: Context) => {
  const url = new URL(request.url);

  // Look for the query parameter, and return if we don't find it
  if (url.searchParams.get("method") !== "transform") {
    return;
  }

  const response = await context.next();
  const text = await response.text();

  const document = new DOMParser().parseFromString(text);
  const h1 = document.querySelector("h1");
  h1.innerHTML = "Thanks for visiting from " + context.geo.city + "!";
  return new Response(document.toString(), response);
};
```


This simple example Netlify Edge Function uses the geolocation information and a DOM parser written for Deno to modify the HTML response.

Vercel Middleware and Edge Functions

- Powered by Cloudflare Workers (though only 18 regions supported)
- Edge functions are built like regular API functions but with the edge runtime defined
- Vercel KV, Vercel Blob, and Vercel Edge Config data offerings
- Middleware runs on every request for applicable paths before hitting the origin while edge functions are an endpoint
- Include built-in geolocation and IP information

Vercel's edge functions run on Cloudflare, but currently support 18 regions. While you don't need to use Next.js to use them, there are a ton of features built into Next.js that integrate directly into Vercel's offerings.

Vercel + LaunchDarkly



Just one final note, if you are interested LaunchDarkly just recently released a Vercel Edge Config integration and Edge SDK. This synchronizes your flag values with the edge config data store, making flag values available at the edge in your apps deployed to Vercel.

Slides and Research



<https://github.com/remotesynth/edge-serverless-preso>

Questions?

- Mastodon: [@remotesynth@mastodon.xyz](https://mastodon.xyz/@remotesynth)
- Email: brinaldi@launchdarkly.com